

Distributed Systems Final Project Design Document

Minqi Ma (JHED: mma17), Jerry Chen (JHED: gchen41)

General Idea

- Each server maintains 5 log files of updates from each server, and one state file for easy recovery from crash.
- Client sends update to the server's public group. Server first saves the update to its log file, and sends the update to other servers.
- Each update has a lamport timestamp of counter, server_index and index. Server maintains a 5*5 matrix of the indices.
- Execute an update
 - Join
 - Do not save in log file
 - Find if the client is in any room previously
 - Inform other servers to remove this client from the old room, and add to the new room
 - Append
 - Create a message with id = counter and server_index of this update.
 - Messages in a room is sorted in the canonical order of the id.
 - Client saves the message id in memory and use it to create like/unlike updates.
 - Like/unlike:
 - If the user has never liked/unliked this message, create a new like node with the username, counter and server_index of this update.
 - If the user has liked/unliked this message, compare saved timestamp and new update's timestamp. Only update the liked status if this update is in a higher canonical order.
 - If any update is out of order, ignore the update.
- **Merge steps**
 - Server with the smallest server index sends a START signal, in case one merge happens in the middle of another.
 - Server sends current participant list for each room to servers group. After merge server informs server-room group of the reconciled participant list.
 - Servers exchange matrices. Server which knows the most about one server's updates will send missing updates for others to catch up with (with coarse-grain flow control).
 - Each server wait until receiving all the missing updates. Sort them in the canonical order and execute one by one.
 - If client sends any updates during merge, put in buffer and execute after merge.
- Server can clean logs in memory if all 5 servers know the updates up to some index.
- Server saves state of data structures for every `FREQ_SAVE` updates received.

Group Architecture

- Server's public group, named "server?". It is used for clients to send updates to the server.

```
#server1#ugrad1
```

- Client's private group, named "#<username>#ugrad?". It is used for server to send messages specifically to this client.
- Server-client group, for each connected server and client pair, named "server?-<username>-ugrad?". It is used to detect client/server connectivity.

```
#server1#ugrad1  
#<username>#ugrad?
```

- Server-room group, for all clients in this room connected to the same server, named "server?-<roomname>". It is used for the server to send room updates

```
#user1#ugrad1  
#user2#ugrad1  
#user1#ugrad2  
.....
```

- Servers group with all 5 servers in it, named "servers". It is used to send messages among servers, and detect servers' connectivity.

```
#server1#ugrad1  
#server2#ugrad2  
#server3#ugrad3  
#server4#ugrad4  
#server5#ugrad5
```

Messages

Chatroom Membership

- JOIN <room_name>: client to server, request to join a room
- ROOMCHANGE <client_name> <old_room> <new_room> <server_index>: server to servers, notify other servers if my client changes room
- PARTICIPANTS_ROOM <user1> <user2>: server to client, inform client of current participants for this room
- PARTICIPANTS_SERVER <room_name> <server_index> <client1> <client2>: server to servers, inform other servers of current participants of a room, used to reconcile on chatroom membership during merge

Updates

- <update> =
 - a <room_name> <username> <content>

- l <room_name> <counter of the liked message> <server_index of the liked message> <username>
- r <room_name> <counter of the unliked message> <server_index of the unliked message> <username>
- UPDATE_CLIENT <update>: client to server, send update to server
- UPDATE_NORMAL <counter> <server_index> <index> <update>: server to servers, stamp the client's update with a lamport timestamp, save it to log file, and send it to other servers
- UPDATE_MERGE <counter> <server_index> <index> <update>: server to servers, reconcile on missing updates during merge
- APPEND <counter of the message> <server_index of the message> <username> <content>: server to client, client will append a new message and save its message id
- LIKES <counter of the message> <server_index of the message> <num_likes>: server to client, client will update number of likes for this message

Others

- MATRIX <25 integers>: server to servers, exchange matrix during merge
- START <5 integers 0/1>: signal to start merging for the current network component
- CONNECT: client to server, request to connect with server
- MESSAGES (<counter> <server_index> <creator> <num_likes> <content>\n)*: server to client, send latest 25 messages of the room
- HISTORY
 - HISTORY <room_name>: client to server, request history of the room
 - HISTORY <creator> <num_likes> <content>: server to client, send the history
- VIEW
 - VIEW: client to server, request membership of each server
 - VIEW <5 integers 0/1>: server to client, 0/1 represents connectivity for each server in the current network component

Data Structures

Client

- struct message
 - *int counter*: message id
 - *int server_index*: message id
 - *char* creator*: username of the creator
 - *char* content*: content of the message
 - *int num_likes*
 - *struct message* next*
- *struct message* messages*: up to 25 messages in this room
- *struct participant*
 - *char* name*
 - *struct participant* next*
- *struct participant* participants*: list of participants in this room
- *char* username*

- *char* room_name*
- *int server_index*: connected server index

Server

- *struct room*
 - *char* room_name*
 - *struct participant* participants[5]*: array of 5 lists of participants from each server
 - *struct message* messages*: list of messages in this room
 - *struct room* next*
- *struct room *rooms*: list of all rooms
- *struct message*
 - *int counter*: message id
 - *int server_index*: message id
 - *char* content*: content of the message
 - *char* creator*: username of the creator
 - *struct like* likes*: list of likes and unlikes of different users
 - *struct message* next*
- *struct like*
 - *char* username*
 - *bool liked*: if this user likes/unlikes the message
 - *int counter*: counter of highest timestamp applied
 - *int server_index*: server_index of highest timestamp applied
 - *struct like *next*
- *struct participant*
 - *char* name*
 - *struct participant* next*
- *int my_counter*
- *int my_index*
- *struct log* logs[5]*: 5 lists of logs from each server
- *struct log*
 - *int counter*
 - *int server_index*
 - *int index*
 - *char* content*
 - *struct log* next*
- *int matrix[5][5]*
- *struct log* buffer*: buffered client updates received in the middle of merging
- *struct log* updates*: updates from log file OR received missing updates during merge
- *bool connected_servers[5]*: current network connectivity
- *int num_updates*: number of updates expected to receive during merge
- *bool received_matrix[5]*: if I have received the matrix from server (i+1) during merge
- *bool merging*: in merging state or not

- *bool received_start*: if received start merging signal, to prevent from receiving updates from previous network change
- *int sent_updates[5]*: highest index of sent missing logs

Algorithm

Client

User command

- u <username>
 - Leave previous server-client, server-room group if exist
 - Clear data structures, like *messages*, *participants* list
 - Create its new private group with the new username
- c <server_index>
 - Record server index; later the client will send updates to this public group
 - Leave previous server-client group, if exists
 - Clear data structures, like *messages*, *participants* list
 - Join new server-client group
 - Send "CONNECT" message to the server
 - Start a timer. If the server is not connected when timer expires, inform user that connection failed.
- j <room_name>
 - Leave previous server-room group, if it exists
 - Join new server-room group
 - Send "JOIN <room_name>" message to the server's public group
- a <content>
 - Send "UPDATE_CLIENT a <room_name> <username> <content>" to the server's public group
- l <line_number>
 - Find the message's id (counter and server_index) in *messages* list
 - Send "UPDATE_CLIENT l <room_name> <counter of the liked message> <server_index of the liked message> <username>" to the server's public group
- r <line_number>
 - Find the message's id (counter and server_index) in *messages* list
 - Send "UPDATE_CLIENT r <room_name> <counter of the unliked message> <server_index of the unliked message> <username>" to the server's public group
- h
 - Send "HISTORY <room_name>" to the server's public group
- v
 - Send "VIEW" to the server's public group

Regular messages

- Receive "MESSAGES ..." from server
 - Construct *messages* list accordingly

- Display
- Receive "PARTICIPANTS_ROOM <user1> <user2> ..." from server-room group
 - Clear *participants* list and reconstruct the list accordingly
 - Display
- Receive "APPEND <counter of the message> <server_index of the message> <username> <content>" from server-room group
 - Append a new message to *messages* list and record its id
 - If list size is larger than 25, remove the first message
 - Display
- Receive "LIKES <counter of the message> <server_index of the message> <num_likes>" from server-room group
 - Search the message id in *messages* list
 - If the message exists, update its number of likes
- Receive "HISTORY <creator> <num_likes> <content>" from server
 - Display
- Receive "VIEW <5 numbers 0/1>" from server
 - Display

Membership messages

- Receive membership change in server-client group
 - JOIN (client itself or server joins the group)
 - If the server joins the group, mark as successfully connected to server
 - DISCONNECT/NETWORK CHANGE (server crashes/daemon crashes)
 - Leave previous server-room group, if exists
 - Clear data structures, like *messages*, *participants* list
 - Leave server-client group
- Receive membership change in server-room group
 - Do nothing; the server will send corresponding updates later

Server

Start: ./server <my_server_index>

- Upon the server starts
 - Join servers group
 - Join its public group "server<my_server_index>"
 - If there is state file
 - Reconstruct data structures from state file
 - Retrieve counter, index and matrix; matrix[my_server_index] vector represents indices of last executed logs to reach this state
 - else
 - initialize empty data structures, *counter* = 0, *index* = 0, matrix is all 0's
 - For every server, if log file exists
 - Check matrix to see if all 5 servers have logs up to some lowest index
 - Loop through each line, if log index <= lowest index, skip

- If log index \leq *matrix[my_server_index][server_index]* (i.e have executed this log to reach the state)
 - Append to *logs[server_index]* list
 - If log index $>$ *matrix[my_server_index][server_index]* (i.e have not executed this log)
 - Insert it to *updates* list in canonical order
- Traverse each log in *updates* list
 - Append the update to *logs[server_index]* list
 - Adopt the counter if higher
 - Update matrix accordingly
 - If the update is from myself, adopt the index
 - Execute the update
- Start receiving messages

Regular messages

- Receive “CONNECT” message in the public group
 - Join server-client group “server?#<client_name>”
- Receive “JOIN <room_name>” message from <client_name> in the public group
 - If in merging state, put it in *buffer* list
 - Search rooms and see if the client is previously in any room
 - Send “ROOMCHANGE <client_name> <old_room> <new_room> <server_index>” to servers group (<old_room> can be “null”)
 - Send up to latest 25 messages of this room to the client’s private group
 - Format: “MESSAGES” followed by messages
 - Each message: <counter> <server_index> <creator> <num_likes> <content>\n
- Receive “ROOMCHANGE <client_name> <old_room> <new_room> <server_index>” in servers group
 - If in merging state, put update it in *buffer*
 - If old_room is not null, remove client from *participants[server_index]* list in the old room
 - if new_room is not null
 - create the new room if the room does not exist
 - add client to *participants[server_index]* list in the new room
 - If I have clients in the affected rooms (i.e *participants[my_server_index]* is not empty)
 - Send new participant list to the affected server-room groups, “PARTICIPANTS_ROOM <user1> <user2> ...”
- Receive “UPDATE_CLIENT <update>” in the public group
 - If in merging state, put update in *buffer* list
 - Increment counter
 - Increment index
 - Stamp the message with counter + server_index + index
 - Write the message to log file

- Append to *logs[my_server_index]* list
- Update *matrix[my_server_index][my_server_index]* to new index
- Send “UPDATE_NORMAL <counter> <my_server_index> <index> <update>” to servers group
- Receive “UPDATE_NORMAL <counter> <server_index> <index> <update>” in servers group
 - If the update is not sent by myself
 - If index is out of order from *matrix[my_server_index][server_index]*, return
 - Write it to log file “*server[my_server_index]-log[server_index].out*”
 - Append it in *logs[server_index]* list
 - Adopt the counter if it is higher
 - Update *matrix[my_server_index][server_index]* to the new index
 - If update is “a <room_name> <username> <content>”
 - Create a new message and insert it to *messages* list of the room in the canonical order of message id
 - If I have clients in the room (i.e *participants[my_server_index]* is not empty)
 - Send “APPEND <counter> <server_index> <username> <content>” to the server-room group
 - If update is “l <room_name> <counter of the liked message> <server_index of the liked message> <username>”
 - If user is the creator of the message, return
 - If there is already a like node in *likes* list with the same username
 - If this update’s timestamp is not higher than saved one, return
 - Adopt counter, server_index of the node to the higher value
 - If the message is previously unliked by this user
 - Update it to liked status
 - If I have clients in the room (i.e *participants[my_server_index]* is not empty)
 - Send “LIKES <message’s counter> <message’s server_index> <num_likes>” to the server-room group
 - If there is no node in *likes* list with the same username
 - Create a new node and append it to *likes* list
 - If I have clients in the room (i.e *participants[my_server_index]* is not empty)
 - Send “LIKES <message’s counter> <message’s server_index> <num_likes>” to the server-room group
 - If update is “r <room_name> <counter of the unliked message> <server_index of the unliked message> <username>”
 - If user is the creator of the message, return
 - If there is already a like node in *likes* list with the same username
 - If this update’s timestamp is not higher than saved one, return

- Adopt counter, *server_index* of the node to the higher value
- If the message is previously liked by this user
 - Update it to unliked status
 - If I have clients in the room (i.e *participants[my_server_index]* is not empty)
 - Send "LIKES <message's counter> <message's server_index> <num_likes>" to the server-room group
- If there is no node in *likes* list with the same username
 - Create a new node and append it to *likes* list
- For every *FREQ_SAVE* updates received
 - Save data structures to state file

(Reconciliation)

- Receive "START <5 integers 0/1>" in servers group
 - If 5 integers do not match with current network connectivity, return
 - *received_start* = true
 - for every room
 - Clear *participants[server_index]* for servers not in the current network component
 - Send *participants[my_server_index]* to the servers group, "PARTICIPANTS_SERVER <room_name> <server_index> <client1> <client2> ..."
 - Send my matrix "MATRIX <25 integers>" to servers group
 - Initialize *received_matrix* array to false for all the connected servers
 - Initialize number of updates expected to receive *num_updates* = 0
- Receive "PARTICIPANTS_SERVER <room_name> <server_index> <client1> <client2> ..." in servers group
 - If *received_start* == false, return
 - Create the room if it does not exist in *rooms* list
 - Clear *participants[server_index]* list in this room and construct new list accordingly
- Receive "MATRIX <25 integers>" in servers group
 - If *received_start* == false, return
 - Adopt all integers if it is higher, except for my vector *matrix[my_server_index]*
 - Mark as received matrix from this server (*received_matrix[server_index]* = true)
 - If *received_matrix* array is all true, (i.e just received all matrices)
 - Clear *updates* list
 - For every room
 - If I have clients in the room (i.e *participants[my_server_index]* is not empty)
 - Send new participant list of this room to server-room group, in the format of "PARTICIPANTS_ROOM <client1> <client2> ..."

- For each server column
 - Clear *logs[server_index]* list up to the lowest index among 5 servers
 - Calculate the number of updates expected to receive and add to *num_updates*
 - If I have the highest log index with the lowest server index
 - Get update in *logs[server_index]* from lowest index+1 to highest index
 - Send INIT_SEND_SIZE updates to servers group in format of "UPDATE_MERGE <counter> <server_index> <index> <update>"
 - Record index of highest sent updates in *sent_updates[server_index]*
 - If *num_updates* == 0 (i.e no updates to merge)
 - Mark as out of merging state
 - If *buffer* list is not empty
 - Execute UPDATE_CLIENT or JOIN or HISTORY or ROOMCHANGE as normal
- Receive "UPDATE_MERGE <counter> <server_index> <index> <update>" in servers group
 - If *received_start* == false, return
 - If *num_updates* == 0, return (i.e it is an update sent in the previous network change)
 - *num_updates*--
 - If update is sent by myself
 - If *sent_updates[server_index]* has not reached the highest log index I have
 - Send one more UPDATE_MERGE to servers group
 - *sent_updates[server_index]*++
 - If I do not have the update (i.e index > matrix[my_server_index][server_index])
 - Insert it in *updates* list in the canonical order
 - If *num_updates* == 0, (i.e just received all missing updates)
 - For every update in *updates* list
 - Write it to log file
 - Append it to *logs[server_index]* list
 - Adopt the counter if it is higher
 - Update *matrix[my_server_index][server_index]* to the new index
 - Execute the update
 - Mark as out of merging state
 - If *buffer* list is not empty
 - Execute UPDATE_CLIENT or JOIN or HISTORY or ROOMCHANGE as normal
- Receive "HISTORY <room_name>" from <client_name> in the public group

- If in merging state, put it in *buffer* list
- for each message in this room
 - Send “HISTORY <creator> <num_likes> <content>” to client’s private group
- Receive “VIEW” from <client_name> in the public group
 - Send “VIEW <5 numbers 0/1>” to the client’s private group

Membership messages

- Receive membership change in server-client group
 - LEAVE/NETWORK CHANGE/DISCONNECT (i.e client reconnects to another server or crashes)
 - Search all rooms and see if the client is previously in any room
 - If the client is previously in a room, send “ROOMCHANGE <client_name> <old_room> <null> <server_index>” to servers group
 - Leave server-client group

(Reconciliation)

- Receive membership change in servers group (i.e servers crash/network partition)
 - Mark as in merging state (i.e *merging* = true)
 - Record network connectivity in *connected_servers* array
 - *received_start* = false
 - If I have the lowest server_index in the current network component
 - Send “START <5 integers 0/1>” representing current network connectivity

File Format

Log

- Name: server<server_index>-log<server_index of logs>.out
e.g server1-log1.out, server1-log2.out
- Line format: <counter> <server_index> <index> <content>
- <content> =
 - a <room_name> <username> <content>
 - l <room_name> <counter of the liked message> <server_index of the liked message> <username>
 - r <room_name> <counter of the unliked message> <server_index of the unliked message> <username>

State

- Name: server<server_index>-state.out
E.g server1-state.out, server2-state.out
- Format (line by line)
 - counter
 - matrix: 5 lines, each line has 5 integers
 - <num_rooms>
 - <room_name> <num_messages>, followed by messages in this room
 - Message:

<counter> <server_index> <creator> <length of *likes* list> <content>

Followed by <like> lines

○ Like:

<username> <counter> <server_index> <liked 0/1>

● e.g state1.txt

10 ← counter

1 2 3 4 5 ← matrix

1 3 4 5 6

6 3 5 6 3

1 3 2 1 1

0 0 0 0 0

2 ← 2 rooms

room1 2

1 1 user1 1 hello everyone ← message line

user2 1 2 1 ← 1 like node

2 1 user2 0 hi there

room2 1

4 2 user3 2 what's up

user1 1 2 0

user2 1 3 1